Cognitive Error–Aware Adaptive Feedback System for Educational Assessments: A Machine Learning Approach to Data-Driven Error Diagnosis and Personalized Learning Interventions

Hemit Rana*, Devesh Sahu*, Jasmeen Kaur*, Nikhil Girsa*, Neftalem Gebremicael*
*Master of Applied Computing (MAC), University of Windsor, Windsor, Canada
Email: rana6c@uwindsor.ca, Sahu21@uwindsor.ca, jasmeen1@uwindsor.ca,
girsa@uwindsor.ca, gebremin@uwindsor.ca

***Abstract*—Online learning platforms have become central to modern education, yet the feedback they provide is typically limited to correctness-based responses that ignore underlying cognitive reasons for student errors. This paper presents a Cognitive Error–Aware Adaptive Feedback System that leverages machine learning to automatically infer latent cognitive error categories from sequences of student responses across assignments and quizzes. The system processes longitudinal assessment data to identify behavioral patterns such as repeated mistakes, inconsistent answers, and improvement trajectories. Unsupervised clustering and sequence analysis models infer error types such as misconceptions, guessing behavior, partial understanding, procedural errors, and careless mistakes. Based on the inferred error category, the system generates targeted adaptive feedback including conceptual reframing for misconceptions, scaffolded hints for guessing behavior, boundary examples for partial understanding, and computational guidance for procedural errors. Unlike rule-based systems requiring extensive expert engineering, this data-driven approach automatically discovers interpretable error representations, enabling scalable personalization. Evaluation will measure error category quality through clustering coherence metrics, learning effectiveness through repeated error reduction and simulated pre/post-assessment deltas, and feedback quality through expert review and relevance assessment. The proposed system is fully software-based, feasible within a semester timeline, and addresses a critical gap in educational technology by shifting artificial intelligence application from performance prediction toward error-aware pedagogical decision-making.**

I. Introduction

A. Problem Motivation and Educational Context

Online learning platforms and digital assessment tools have become central to modern education, particularly in large-scale courses and remote learning environments. While these systems efficiently evaluate student performance, the feedback they provide is predominantly limited to correctness-based responses. Students typically receive notification that their answer is correct or incorrect, sometimes accompanied by generic explanations. This uniform feedback model rests on a critical and largely unexamined assumption: that all incorrect answers stem from equivalent knowledge gaps and warrant identical interventions.

In reality, students make errors for fundamentally different cognitive reasons. Consider a typical scenario in an introductory physics course where three students answer a question about momentum incorrectly. Student A guesses randomly without attempting calculation. Student B has constructed a systematic misconception, confusing momentum with velocity, and will likely apply this flawed reasoning consistently. Student C applies the correct formula but makes a computational arithmetic error. Current assessment systems treat all three identically, providing the same generic explanation. Research in cognitive science and learning sciences demonstrates that this one-size-fits-all approach significantly reduces learning effectiveness and leads to repeated mistakes.

B. Why This Problem Matters

The learning sciences literature consistently demonstrates that feedback tailored to the specific source of error substantially improves learning outcomes and reduces time-to-mastery. Educational psychology establishes that students hold systematic misconceptions that persist across multiple questions, not merely temporary gaps in knowledge. Large-scale assessments generate enormous volumes of response data but fail to extract actionable insights about error semantics. This represents a critical missed opportunity for artificial intelligence application. While grading automation has largely been solved, the deeper problem of understanding learner behavior and enabling effective intervention design remains largely unexplored in practical systems, despite strong theoretical foundations and evidence for its effectiveness.

II. Core Project Idea and System Overview

A. System Architecture

We propose a Cognitive Error–Aware Adaptive Feedback System that processes longitudinal student assessment data to automatically discover interpretable error categories and generate personalized feedback. The system operates through five integrated stages.

First, the data collection stage gathers student response sequences including answers, correctness labels, attempt counts, response times, and temporal ordering across multiple assignments and quizzes. This longitudinal perspective is critical for detecting persistent patterns distinguishing misconceptions from transient errors.

Second, the feature extraction stage computes behavioral features capturing response patterns. These include repeated mistakes on structurally similar problems, indicating persistent cognitive gaps; inconsistency in responses to equivalent questions, revealing unstable understanding; improvement trajectories and plateaus showing learning progress or stagnation; response time anomalies that may indicate careless errors versus deliberate misunderstanding; and attempt count patterns distinguishing lucky guesses from systematic problem-solving.

Third, the error category inference stage applies unsupervised learning techniques to identify latent error types. Clustering algorithms group students by response patterns, and sequence analysis models detect temporal dependencies. Rather than using predetermined error categories, the system discovers meaningful archetypes from data, enabling generalization across domains and learner populations.

Fourth, the adaptive feedback generation stage maps inferred error categories to targeted interventions. For misconception errors characterized by systematic incorrect reasoning, feedback provides conceptual reframing and contrastive examples showing correct versus incorrect reasoning. For guessing behavior indicated by random answers without engagement, feedback employs scaffolded hints guiding the problem-solving process. For partial understanding where students succeed in some contexts but not others, feedback provides boundary examples illustrating principle limits. For procedural errors where conceptual understanding is present but execution falters, feedback offers step-by-step computational guidance. For careless errors characterized by capability with temporary lapses, feedback employs checklist strategies and attention-focusing techniques.

Fifth, the synthesis of follow-up questions stage generates contextually appropriate follow-up questions designed to reinforce correct reasoning, probe understanding depth, and guide continued learning.

B. Core Insight and Innovation

The fundamental innovation is shifting the application of artificial intelligence in educational assessment from performance prediction and automated grading toward error-aware pedagogical decision-making. Most existing assessment systems assume that incorrect answers uniformly indicate knowledge deficiency. Rule-based approaches to feedback routing require extensive expert input and fail to generalize. Standard machine learning approaches often remain in the performance prediction space, predicting future correctness without characterizing the nature or source of errors in interpretable pedagogical terms.

This project reframes adaptive feedback generation as a data-driven error discovery problem. Rather than predicting whether students will improve or which students are at risk, the system asks a distinct question: What type of cognitive gap or misconception underlies this error, and what intervention specifically addresses that gap? By modeling error semantics rather than merely correctness, the system enables interventions tailored to actual learning needs, a capability that remains underexplored despite strong evidence for its pedagogical effectiveness.

III. AI Novelty and Technical Contribution

A. Challenging the Uniformity Assumption

The prevailing paradigm in educational assessment systems treats incorrect answers as equivalent indicators of knowledge deficiency. Existing educational AI has made substantial progress on performance prediction and automated grading, tasks that have largely matured and are now commoditized. However, these systems have largely neglected the pedagogically critical problem of error semantics: understanding not merely whether a student will succeed, but why they erred and what type of intervention addresses their specific cognitive gap.

Rule-based systems represent the current standard for adaptive feedback. These require extensive expert engineering to define error categories and craft feedback templates. They fail to generalize across domains and learner populations and require complete redesign when moving to new subjects or assessment contexts. Standard machine learning approaches have focused predominantly on correctness prediction and knowledge level classification, which, while valuable for institutional efficiency, provide no actionable understanding of error sources.

B. Research Gap and Contribution

The gap we address is the lack of practical systems that automatically learn error-type-specific patterns from data and couple this inference with pedagogically-informed adaptive feedback. Recent literature in educational data mining emphasizes the importance of temporal patterns and behavioral trajectories but lacks operational systems that implement this insight at scale. Our contribution is threefold: first, we operationalize error discovery as an unsupervised learning problem on response sequences and behavioral features; second, we develop interpretable error categories that align with learning science theory; third, we couple error inference with templated adaptive feedback generation that makes pedagogical sense.

This differs fundamentally from prior work in several ways. Knowledge tracing systems like Deep Knowledge Tracing predict the probability of correctness but do not categorize error types. Hint generation systems like those by Stamper et al. [1] predict when hints are needed but do not diagnose why errors occur. Error detection systems focus on whether an error occurred, not what type of error or what intervention addresses it. Our approach integrates sequence analysis, clustering, and supervised classification specifically to solve the error diagnosis problem in an interpretable, actionable way.

IV. Technical Approach and Implementation

A. System Architecture

The system architecture consists of five interconnected modules. The Data Pipeline Module loads student response sequences from educational datasets and performs normalization and feature engineering. The Feature Extraction Module computes behavioral features capturing response patterns including error consistency, improvement trajectories, response time variability, and temporal clustering. The Error Category Inference Module applies clustering and sequence analysis to discover latent error types. The Adaptive Feedback Synthesis Module maps error categories to targeted feedback templates. The Evaluation Framework Module measures system effectiveness through multiple metrics.

B. Technical Implementation Details

We implement the system in Python using scikit-learn for machine learning, PyTorch for potential neural sequence models, pandas and numpy for data processing, and SQLite for data storage. The database system provides declarative abstraction through SQL queries, logical-physical data independence enabling optimization, and indexing on student IDs and question IDs for efficient retrieval of response sequences.

The Feature Extraction Module computes approximately twenty behavioral features per student-question pair. These include error counts on specific questions and similar questions, correction counts measuring recovery from errors, attempt counts before correctness, error consistency scores quantifying similarity of incorrect answers, response time variability indicating careless versus deliberate errors, improvement trends measured as accuracy slope over time, and cross-question error rates on structurally similar problems.

The Error Category Inference Module implements multiple complementary approaches. K-means clustering operates on feature vectors to identify natural groupings of student response patterns. Hierarchical clustering provides dendrograms revealing relationship structure between error types. Hidden Markov Models capture temporal dependencies in response sequences, distinguishing transient mistakes from persistent misconceptions. Silhouette scores and Davies-Bouldin indices validate cluster quality and interpretability. Qualitative validation with domain experts ensures discovered categories align with learning science principles.

C. Algorithm Selection Rationale

Unsupervised clustering was selected because it discovers error patterns without requiring labeled error categories, a labor-intensive ground truth that would limit scalability. The results are inherently interpretable, enabling educational experts to validate and refine discovered categories. Sequence analysis methods capture temporal dependencies, distinguishing persistent misconceptions maintained across multiple attempts from transient mistakes that students self-correct. Feature engineering draws on educational research identifying behavioral indicators of different error types, combining domain knowledge with data-driven discovery to balance interpretability and predictive power.

V. Data and Evaluation Plan

A. Datasets and Data Characteristics

The primary dataset is ASSISTments, a publicly available educational data repository containing fifty thousand or more student responses across multiple courses with rich metadata including student IDs, problem IDs, correctness labels, attempt counts, and response times. This dataset covers diverse domains including mathematics and science, enabling evaluation of system generalization.

Secondary datasets include NAEP (National Assessment of Educational Progress) assessment data, available through appropriate institutional access, which provides large-scale educational assessment data enabling cross-domain and cross-population generalization testing. These datasets are characteristic of typical online learning platforms and represent realistic assessment scenarios.

B. Evaluation Methodology

Evaluation employs multiple complementary approaches addressing different aspects of system effectiveness. For error category quality, unsupervised metrics including silhouette coefficient measure cluster cohesion and separation, Davies-Bouldin index measures cluster compactness, and qualitative interpretability assessment determines whether educators can understand and validate discovered categories. For learning effectiveness, repeated error reduction measures whether students make fewer errors on similar problems after receiving error-specific feedback, synthetic pre/post-assessment deltas simulate learning gains by applying a Bayesian knowledge tracing model before and after intervention, and baseline comparison contrasts error-aware feedback against correctness-only feedback through learning curve simulation. For feedback quality, domain expert review assesses pedagogical soundness, relevance scores measure how well generated feedback addresses diagnosed errors, and actionability analysis determines whether feedback enables students to understand and apply corrections.

C. Experimental Design

Experiment one addresses error category discovery. We apply clustering algorithms to feature-extracted student data, evaluate cluster quality through standard metrics, and validate interpretability through manual review and educator feedback. This experiment addresses the fundamental question of whether meaningful error patterns exist in student response data.

Experiment two evaluates feedback effectiveness through simulation. We generate error-aware feedback for test-set students, employ a Bayesian knowledge tracing model to simulate learning trajectories with and without error-specific feedback, and compare learning curves. This addresses the core hypothesis that error-aware feedback improves learning more than generic correctness feedback.

Experiment three performs ablation analysis. We gradually remove system components including sequence information, temporal features, and clustering steps, measuring performance degradation. This isolates the contribution of different components and identifies critical system elements.

Experiment four tests cross-domain generalization. We train models on one subject or domain, evaluate error discovery on held-out domains, and assess transfer learning effectiveness. This addresses system robustness and practical applicability across diverse educational contexts.

VI. Expected Outcomes and Impact

A. System Deliverables

The project will produce four primary deliverables. First, a set of interpretable error categories discovered from student response data, characterized and exemplified for educational practitioners. Second, adaptive feedback generation templates mapped to error types, enabling personalized interventions addressing specific cognitive gaps. Third, per-student diagnostic reports showing error profiles and evolution over time, providing teachers actionable insights into learner needs. Fourth, comprehensive evaluation metrics and baseline comparisons quantifying system effectiveness.

B. Educational Impact

This system addresses fundamental challenges in educational technology. It improves learning efficiency by enabling targeted feedback addressing specific misconceptions rather than generic explanations. It provides teachers deeper insight into student cognition through error categorization, informing instructional design and intervention sequencing. It enables personalization at scale by automating error diagnosis, removing the teacher workload barrier that typically limits personalized instruction. It demonstrates the feasibility of using machine learning for pedagogically-informed decision support rather than purely predictive applications.

C. Broader Scientific Contribution

This project contributes to a paradigm shift in educational AI. Rather than focusing exclusively on performance prediction and automated grading, it demonstrates the value of AI for understanding and supporting learning. The approach is domain-agnostic and extensible to diverse educational contexts from K-12 assessment to massive open online courses. By showing that interpretable error categories can be automatically discovered from behavioral data, the work provides a foundation for future systems coupling error diagnosis with adaptive intervention.

VII. Project Timeline and Feasibility

The implementation timeline spans fourteen weeks. Weeks one through two focus on setup and data preparation, establishing data pipelines and feature extraction infrastructure. Weeks three through five focus on error category discovery, implementing clustering algorithms and analyzing results. Weeks six through eight focus on feedback synthesis, developing the feedback generation module and template library. Weeks nine through twelve focus on evaluation and experimentation, executing the four experiments and analyzing results. Weeks thirteen through fourteen focus on documentation and finalization.

The project is feasible within this timeline because all components rely on well-established machine learning libraries with mature implementations. Publicly available datasets eliminate data collection overhead. Incremental implementation allows for early validation and iteration. The fourteen-week timeline is conservative, providing buffer for unexpected challenges or deeper investigation of promising directions.

VIII. Related Work and Literature Positioning

A. Closest Related Work

Stamper et al. [1] developed an automated hint generation system using knowledge tracing to predict which students needed help. This foundational work demonstrates the value of data-driven approaches to tutoring system intervention. However, it focuses on predicting correctness and identifying which students require assistance rather than categorizing the type of error underlying incorrect performance. Our work extends this by explicitly modeling error semantics, distinguishing misconceptions from guessing versus partial understanding, and tailoring feedback based on inferred error category rather than only on predicted knowledge deficiency.

Koedinger and Aleven [2] studied the assistance dilemma, investigating whether more specific hints or more general guidance better supports learning. This work provides pedagogical foundations for our feedback intervention design. However, their research required manual specification of when different hint types would be appropriate. Our approach automates the discovery of when and to whom specific feedback types should be directed based on inferred error patterns.

B. Knowledge Tracing and Sequence Modeling

Piech et al. [3] introduced Deep Knowledge Tracing, demonstrating that recurrent neural networks effectively model temporal student knowledge by processing sequential assessment data. Their work established that longitudinal response sequences contain rich signals for understanding learner behavior. We adopt this insight about sequence value but apply it specifically to error pattern discovery rather than knowledge level prediction. Recent systematic reviews by Minn [4] document the shift from purely predictive modeling toward interpretable and adaptive learning support, positioning our work within this emerging paradigm.

C. Error Analysis and Misconception Research

Foundational educational psychology research by VanLehn [5] and Ohlsson and Langley [6] established that incorrect answers often stem from systematic misconceptions rather than random mistakes. Learners construct coherent but incorrect mental models that persist across contexts. Building on this foundation, projects like SMART Tutoring Systems [7] attempted to codify misconceptions through manual knowledge engineering. Our contribution is operationalizing misconception discovery as an automated, data-driven process that scales beyond hand-coded systems.

D. Adaptive Feedback Systems

Metcalf and Buss [8] review evidence that feedback tailored to error type rather than uniform correctness feedback significantly improves learning outcomes. This establishes the pedagogical value of error-aware feedback. However, their review documents that most existing systems determine feedback routing through manual rules requiring extensive domain expertise. Our approach maintains the pedagogical benefits of error-aware feedback while eliminating manual engineering through automatic error category inference.

Recent work by Liang et al. [9] emphasizes capturing temporal patterns and error trajectories in student modeling, noting this remains underutilized in many systems. Our project operationalizes this emphasis by leveraging sequence clustering and feature extraction from response patterns, treating error discovery as an unsupervised learning problem followed by supervised mapping to adaptive interventions.

IX. Conclusion

This project addresses a concrete and significant educational challenge: the ineffectiveness of generic correctness-only feedback in online assessment systems. By applying machine learning to automatically discover cognitive error categories and generate targeted interventions, the system enables personalized learning support at scale. Unlike existing systems that focus on performance prediction or require manual rule engineering, this approach brings artificial intelligence to bear on the pedagogically critical problem of error diagnosis and adaptive intervention design.

The technical approach is sound, leveraging well-established machine learning techniques applied to a novel problem in educational contexts. The evaluation plan is comprehensive, addressing error category quality, learning effectiveness, feedback quality, and system robustness. Implementation within a semester is realistic using standard tools and publicly available datasets. The expected outcomes include both practical value, through personalized feedback improving learning outcomes, and research contributions, through demonstrating the feasibility of data-driven error diagnosis in education.

This project demonstrates how applied AI can enhance educational assessment systems by enabling deeper understanding of learner behavior and more effective feedback strategies. By improving feedback quality rather than merely automating grading, the system addresses a real-world educational challenge and contributes to the emerging paradigm of AI for learning support and personalization rather than purely administrative automation.

References

[1] J. Stamper, T. Murray, R. Madachy, and K. Malzahn, "Towards closing the loop: Automated data-driven hint generation for tutoring systems," in *Proceedings of the 4th International Conference on Educational Data Mining*, 2011, pp. 139–148.

[2] K. R. Koedinger and V. Aleven, "Exploring the assistance dilemma: Balancing support and learning," *Journal of Learning Sciences*, vol. 16, no. 1, pp. 1–39, 2007.

[3] C. Piech et al., "Deep knowledge tracing," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 28, 2015, pp. 505–513.

[4] A. Y. Minn, "Artificial intelligence and education: A systematic review and perspective," *International Journal of Artificial Intelligence in Education*, vol. 31, no. 1, pp. 78–112, 2021.

[5] K. VanLehn, "Rule learning by indirect instruction," in *Handbook of Classroom Assessment*, 1999, pp. 147–201.

[6] S. Ohlsson and P. Langley, "Identifying solution paths in cognitive diagnosis," *AAAI Technical Report FS-85-01*, 1985.

[7] K. R. Koedinger, J. C. Stamper, M. D. Baker, and J. R. Anderson, "Smart tutoring systems: Learning with a smarter tutor," in *Advances in Intelligent Tutoring Systems*, Springer, 2012, pp. 123–156.

[8] B. Metcalf and R. R. Buss, "Automated feedback in practice: Designing and evaluating adaptive feedback for learning," *Educational Technology Research and Development*, vol. 68, no. 4, pp. 1823–1849, 2020.

[9] X. Liang, Y. Yang, and B. du Boulay, "Student modeling in adaptive learning: A systematic review and meta-analysis," *Journal of Educational Technology & Society*, vol. 25, no. 2, pp. 52–78, 2022.